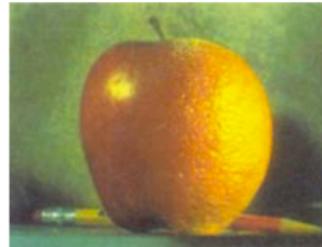


2. Image Formation



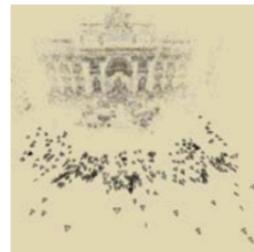
3. Image Processing



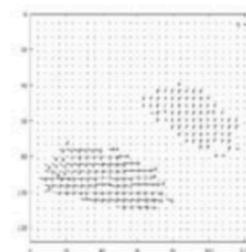
4. Features



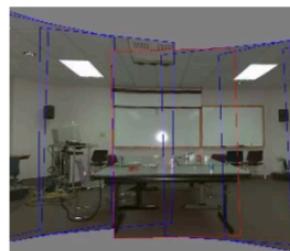
5. Segmentation



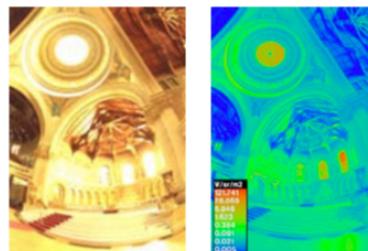
6-7. Structure from Motion



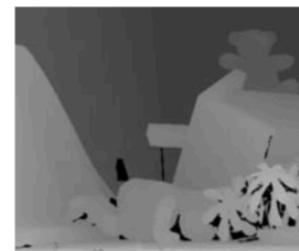
8. Motion



9. Stitching



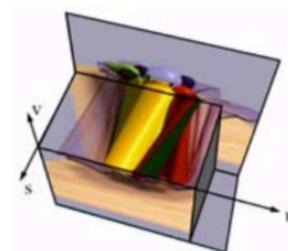
10. Computational Photography



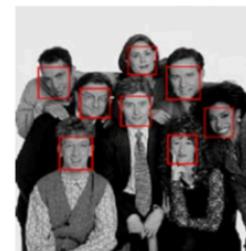
11. Stereo



12. 3D Shape

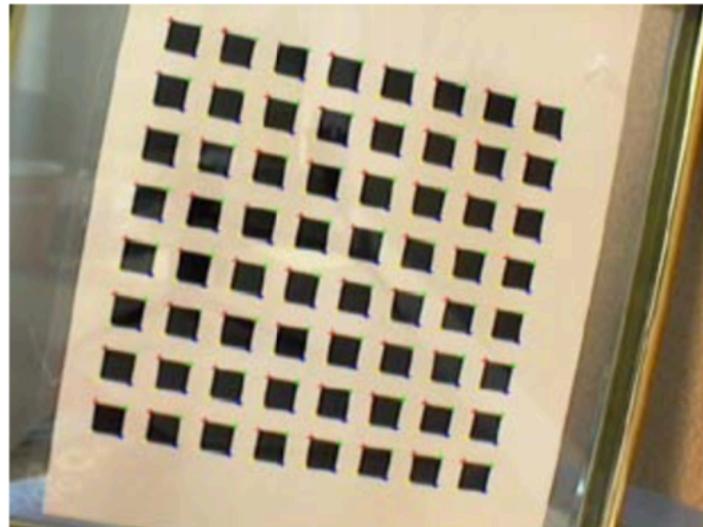
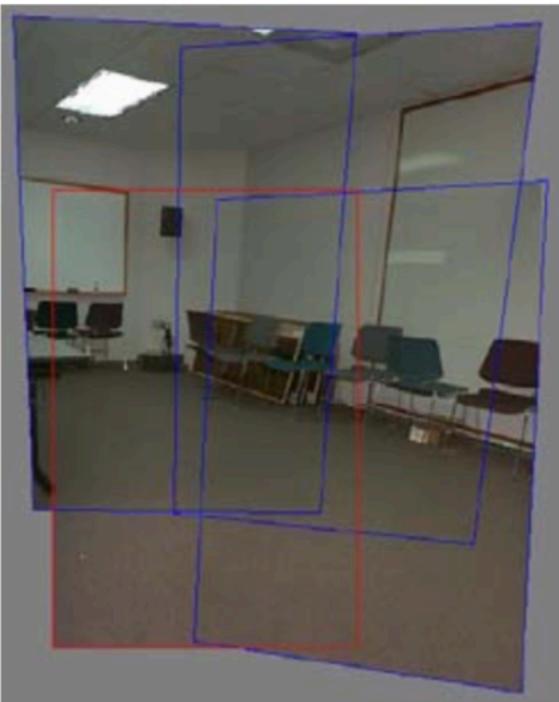


13. Image-based Rendering



14. Recognition

# Feature-based Image Alignment



- Geometric image registration
  - 2D or 3D transforms between them
  - Special cases: pose estimation, calibration

# 2D Alignment



- 3 photos
- Translational model

# 2D Alignment



- Input:
  - A set of matches  $\{(x_i, x'_i)\}$
  - A parametric model  $f(x; p)$
- Output:
  - Best model  $p^*$
- How?

# 2D translation estimation



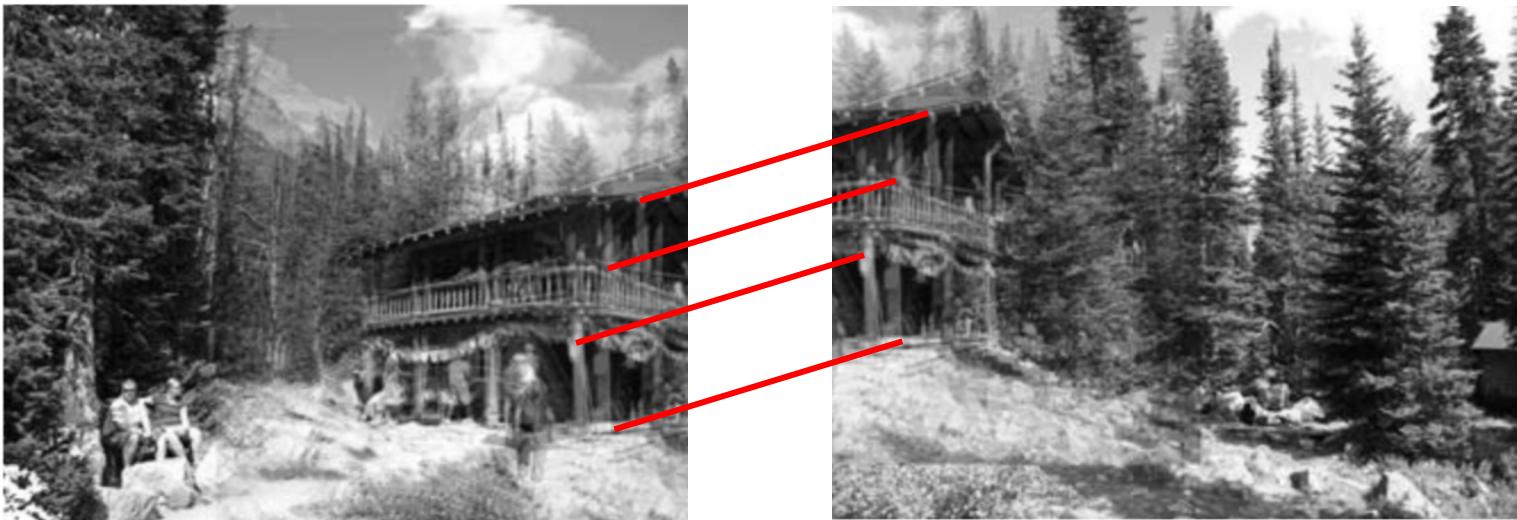
- Input:
  - Set of matches  $\{(x_1, x'_1), (x_2, x'_2), (x_3, x'_3), (x_4, x'_4)\}$
  - Parametric model:  $f(x; t) = x + t$
  - Parameters  $p == t$ , location of origin of A in B
- Output:
  - Best model  $p^*$

# 2D translation estimation



- Input:
  - Set of matches  $\{(x_1, x'_1), (x_2, x'_2), (x_3, x'_3), (x_4, x'_4)\}$
  - Parametric model:  $f(x; t) = x + t$
  - Parameters  $p == t$ , location of origin of A in B
- Question for class:
  - What is your best guess for model  $p^* ??$

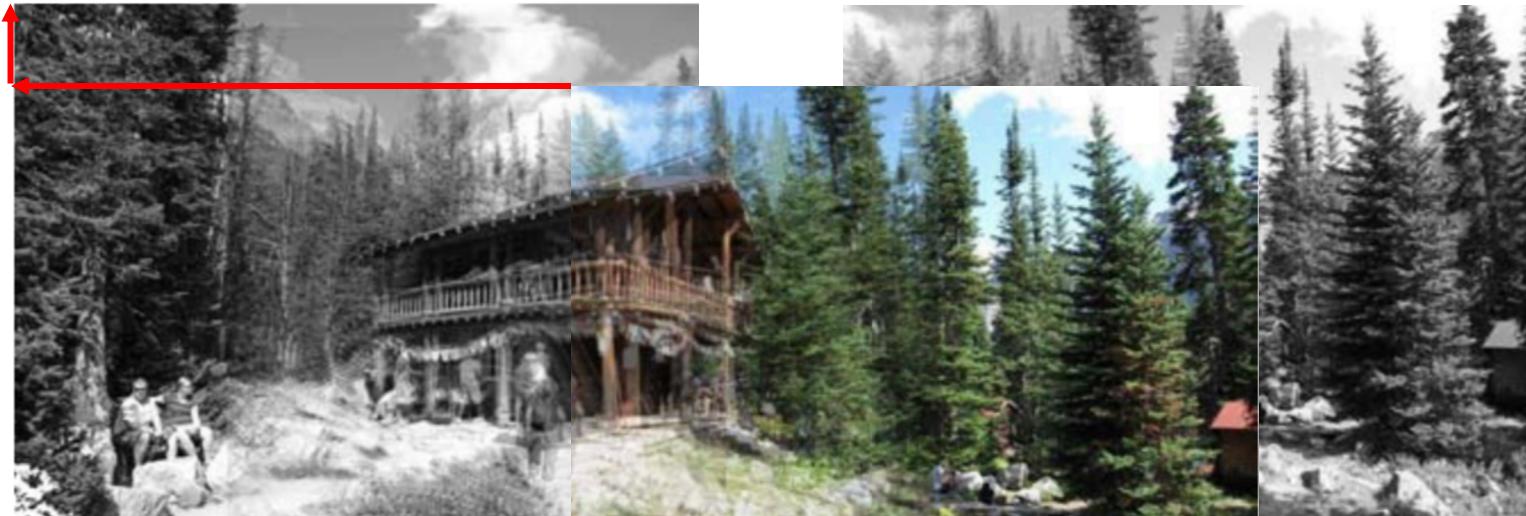
# 2D translation estimation



- How?
  - One correspondence  $x_1 = [600, 150]$ ,  $x_1' = [50, 50]$
  - Parametric model:  $x' = f(x; t) = x + t$

# 2D translation estimation

[-550, -100]



- How?
  - One correspondence  $x_1 = [600, 150]$ ,  $x_1' = [50, 50]$
  - Parametric model:  $x' = f(x; t) = x + t$   
 $\Rightarrow t = x' - x$
  - $\Rightarrow t = [50-600, 40-150] = [-550, -100]$

# 2D translation via least-squares



- How?
  - A set of matches  $\{(x_i, x'_i)\}$
  - Parametric model:  $f(x; t) = x + t$
  - **Minimize sum of squared residuals:**

$$E_{\text{LS}} = \sum_i \|r_i\|^2 = \sum_i \|f(x_i; p) - x'_i\|^2.$$

# How to solve?

In many cases, parametric model is linear:

$$f(x; p) = x + J(x)p$$

$$\Delta x = x' - x = J(x)p$$

$$E_{LS} = \sum_i \|J(x)p + x - x'_i\|^2 = \sum_i \|J(x_i)p - \Delta x_i\|^2$$

Differentiate and set to 0:

$$2 \sum_i J^T(x_i) (J(x_i)p - \Delta x_i) = 0$$

*Normal equations* —  $\left[ \sum_i J^T(x_i) J(x_i) \right] p = \sum_i J^T(x_i) \Delta x_i$

$$Ap = b$$

*Hessian*

$$p^* = A^{-1}b$$

# Linear models menagerie

Transform	Matrix	Parameters $p$	Jacobian $J$
translation	$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix}$	$(t_x, t_y)$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
Euclidean	$\begin{bmatrix} c_\theta & -s_\theta & t_x \\ s_\theta & c_\theta & t_y \end{bmatrix}$	$(t_x, t_y, \theta)$	$\begin{bmatrix} 1 & 0 & -s_\theta x - c_\theta y \\ 0 & 1 & c_\theta x - s_\theta y \end{bmatrix}$
similarity	$\begin{bmatrix} 1 + a & -b & t_x \\ b & 1 + a & t_y \end{bmatrix}$	$(t_x, t_y, a, b)$	$\begin{bmatrix} 1 & 0 & x & -y \\ 0 & 1 & y & x \end{bmatrix}$
affine	$\begin{bmatrix} 1 + a_{00} & a_{01} & t_x \\ a_{10} & 1 + a_{11} & t_y \end{bmatrix}$	$(t_x, t_y, a_{00}, a_{01}, a_{10}, a_{11})$	$\begin{bmatrix} 1 & 0 & x & y & 0 & 0 \\ 0 & 1 & 0 & 0 & x & y \end{bmatrix}$

- All the simple 2D models are linear!
- Exception: perspective transform

# 2D translation via least-squares



For translation:  $J = I$  and normal equations are particularly simple:

$$\left[ \sum_i I^T I \right] p = \sum_i \Delta x_i$$

$$p^* = \frac{1}{n} \sum_i \Delta x_i$$

In other words: just **average** the “flow vectors”  $\Delta x = x' - x$

# Oops I lied !!! Euclidean is not linear!

Transform	Matrix	Parameters $p$	Jacobian $J$
translation	$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix}$	$(t_x, t_y)$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
Euclidean	$\begin{bmatrix} c_\theta & -s_\theta & t_x \\ s_\theta & c_\theta & t_y \end{bmatrix}$	$(t_x, t_y, \theta)$	$\begin{bmatrix} 1 & 0 & -s_\theta x - c_\theta y \\ 0 & 1 & c_\theta x - s_\theta y \end{bmatrix}$
similarity	$\begin{bmatrix} 1 + a & -b & t_x \\ b & 1 + a & t_y \end{bmatrix}$	$(t_x, t_y, a, b)$	$\begin{bmatrix} 1 & 0 & x & -y \\ 0 & 1 & y & x \end{bmatrix}$
affine	$\begin{bmatrix} 1 + a_{00} & a_{01} & t_x \\ a_{10} & 1 + a_{11} & t_y \end{bmatrix}$	$(t_x, t_y, a_{00}, a_{01}, a_{10}, a_{11})$	$\begin{bmatrix} 1 & 0 & x & y & 0 & 0 \\ 0 & 1 & 0 & 0 & x & y \end{bmatrix}$

- ~~All the simple 2D models are linear!~~
- Euclidean Jacobians are a function of  $\theta$ !

# Nonlinear Least Squares

$$E_{NLS} = \sum_i \|f(x_i; p) - x'_i\|^2$$

Linearize around a current guess  $p$ :

$$f(x; p + \Delta p) = f(x; p) + J(x; p)\Delta p$$

$$r = x' - f(x; p) = J(x; p)\Delta p$$

$$E_{NLS} = \sum_i \|f(x; p) + J(x; p)\Delta p - x'_i\|^2 = \sum_i \|J(x; p)\Delta p - r_i\|^2$$

Differentiate and set to 0:

$$2 \sum_i J^T(x_i; p) (J(x_i; p)\Delta p - r_i) = 0$$

$$\left[ \sum_i J^T(x_i; p) J(x_i; p) \right] \Delta p = \sum_i J^T(x_i; p) r_i$$

$$A\Delta p = b$$

$$\Delta p* = A^{-1}b$$

# Projective/H

- Jacobians a bit harder
- Parameterization:

$$\begin{bmatrix} 1 + h_{00} & h_{01} & h_{02} \\ h_{10} & 1 + h_{11} & h_{12} \\ h_{20} & h_{21} & 1 \end{bmatrix}$$

Image credit Graphics Mill  
(educational Use)

$$(h_{00}, h_{01}, \dots, h_{21})$$

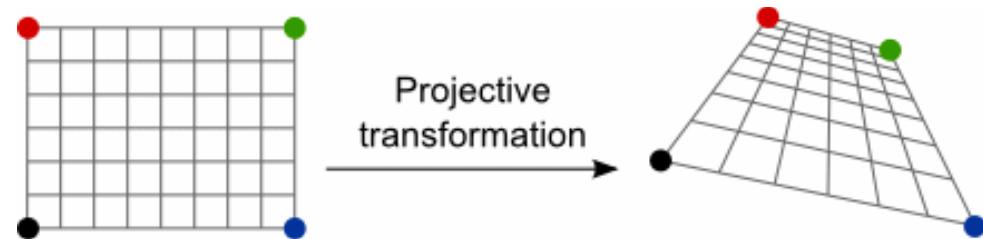
- $\mathbf{x}' = f(\mathbf{x}, \mathbf{p})$ :

$$x' = \frac{(1 + h_{00})x + h_{01}y + h_{02}}{h_{20}x + h_{21}y + 1} \quad \text{and} \quad y' = \frac{h_{10}x + (1 + h_{11})y + h_{12}}{h_{20}x + h_{21}y + 1}.$$

- And Jacobian:

$$\mathbf{J} = \frac{\partial \mathbf{f}}{\partial \mathbf{p}} = \frac{1}{D} \begin{bmatrix} x & y & 1 & 0 & 0 & 0 & -x'x & -x'y \\ 0 & 0 & 0 & x & y & 1 & -y'x & -y'y \end{bmatrix}$$

$$D = h_{20}x + h_{21}y + 1$$



# Closed Form H

- Taking  $x' = f(x, p)$ :

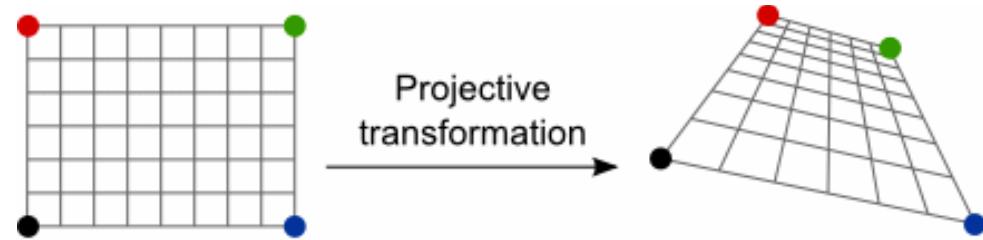


Image credit Graphics Mill

$$x' = \frac{(1 + h_{00})x + h_{01}y + h_{02}}{h_{20}x + h_{21}y + 1} \quad \text{and} \quad y' = \frac{h_{10}x + (1 + h_{11})y + h_{12}}{h_{20}x + h_{21}y + 1}.$$

- Divide both sides by  $D = h_{20}x + h_{21}y + 1$ :

$$\begin{bmatrix} \hat{x}' - x \\ \hat{y}' - y \end{bmatrix} = \begin{bmatrix} x & y & 1 & 0 & 0 & 0 & -\hat{x}'x & -\hat{x}'y \\ 0 & 0 & 0 & x & y & 1 & -\hat{y}'x & -\hat{y}'y \end{bmatrix} \begin{bmatrix} h_{00} \\ \vdots \\ h_{21} \end{bmatrix}$$

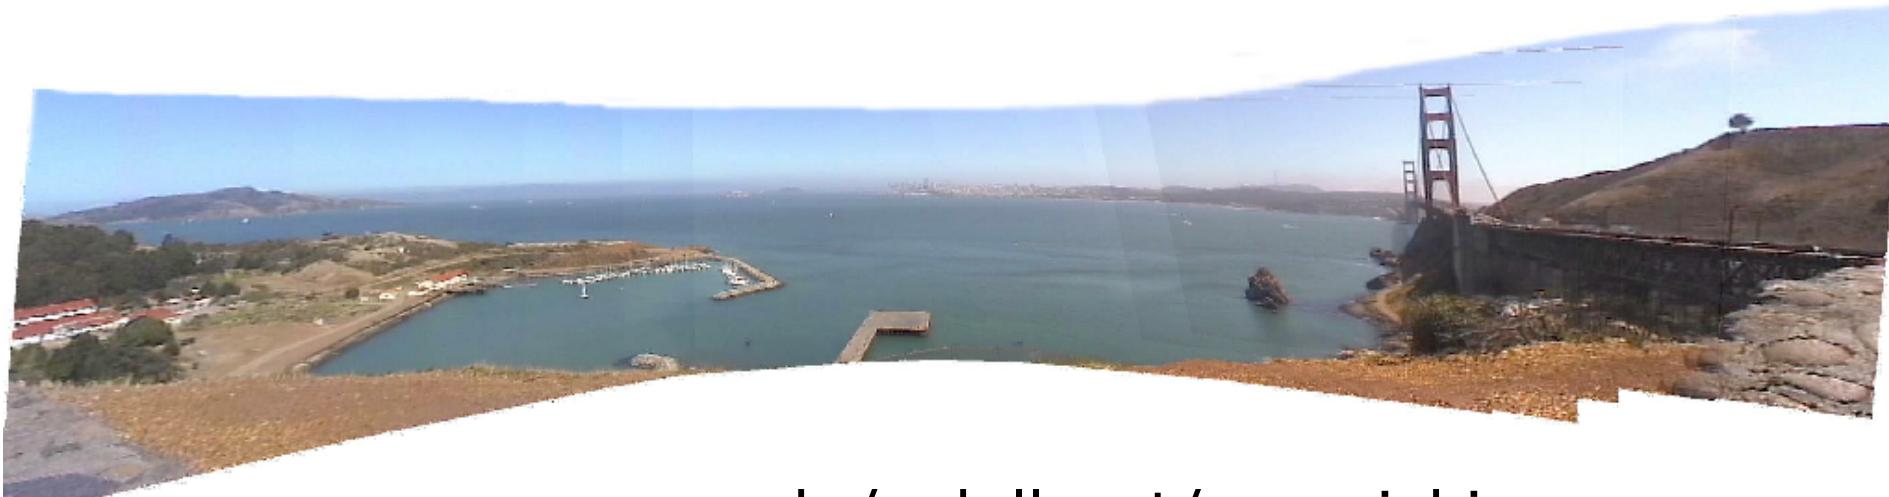
- 4 matches  $\Rightarrow$  system of 8 linear equations

# RANSAC

# Motivation

- Estimating motion models
- Typically: points in two images
- Candidates:
  - Translation
  - Homography
  - Fundamental matrix

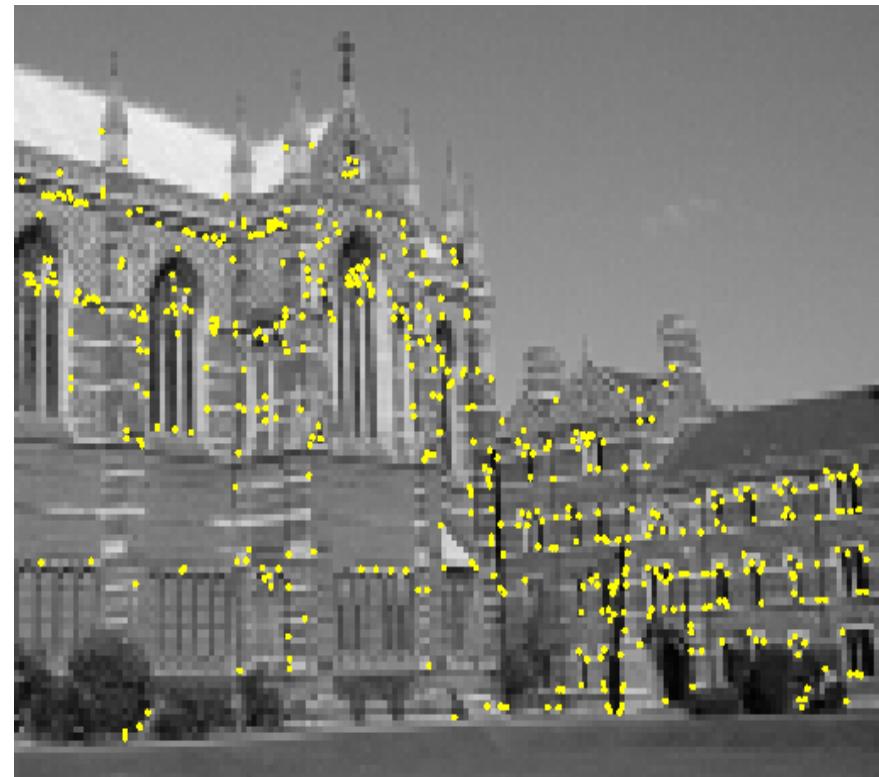
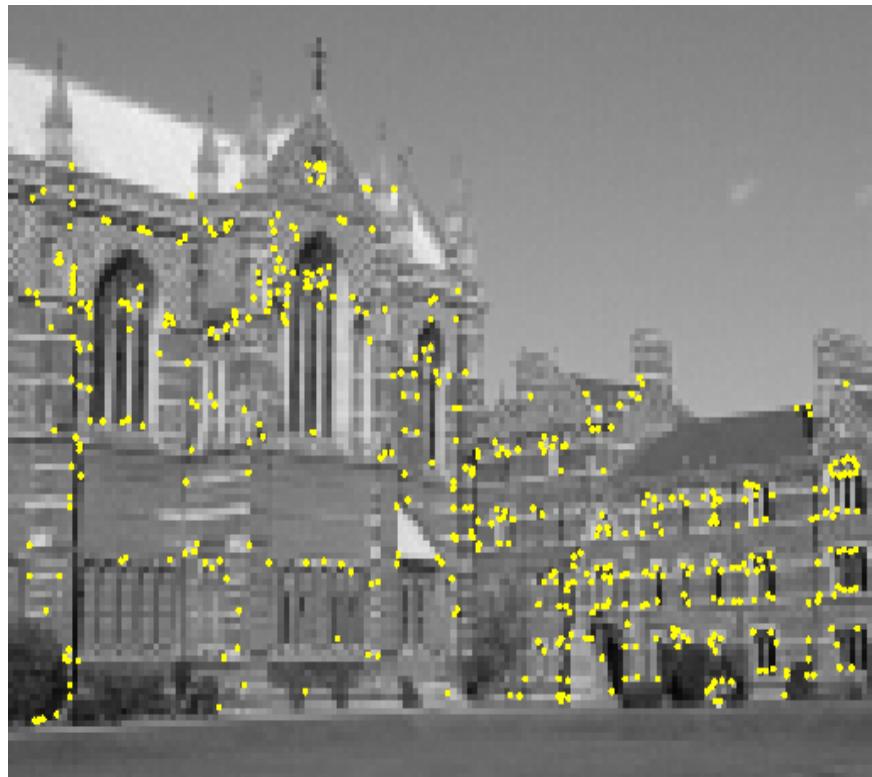
# Mosaicking: Homography



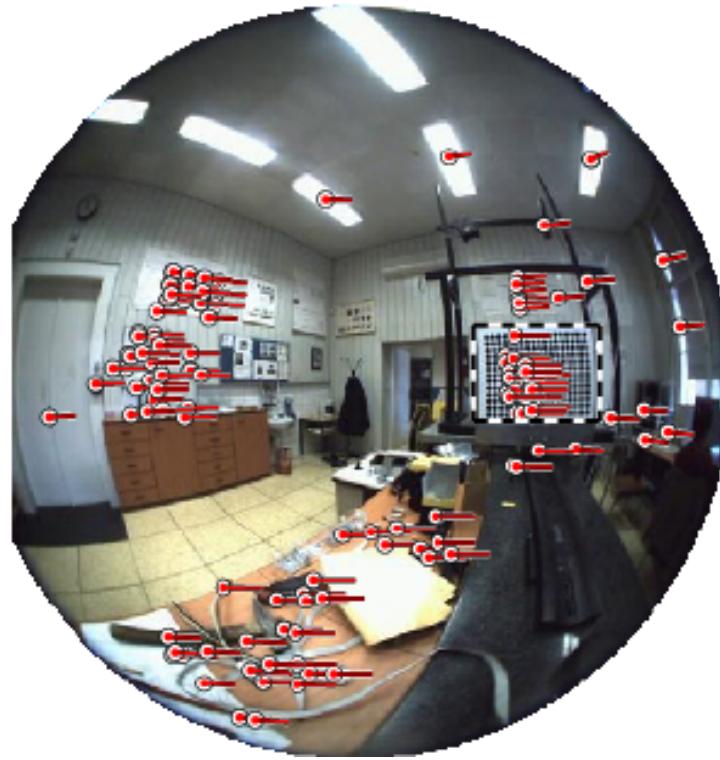
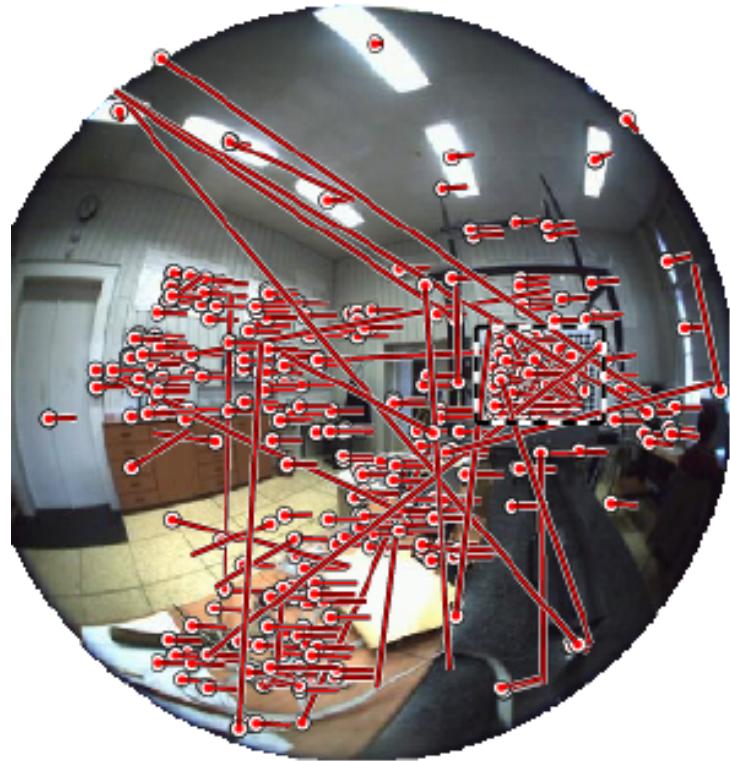
[www.cs.cmu.edu/~dellaert/mosaicking](http://www.cs.cmu.edu/~dellaert/mosaicking)

Frank Dellaert Fall 2020

# Two-view geometry (next lecture)



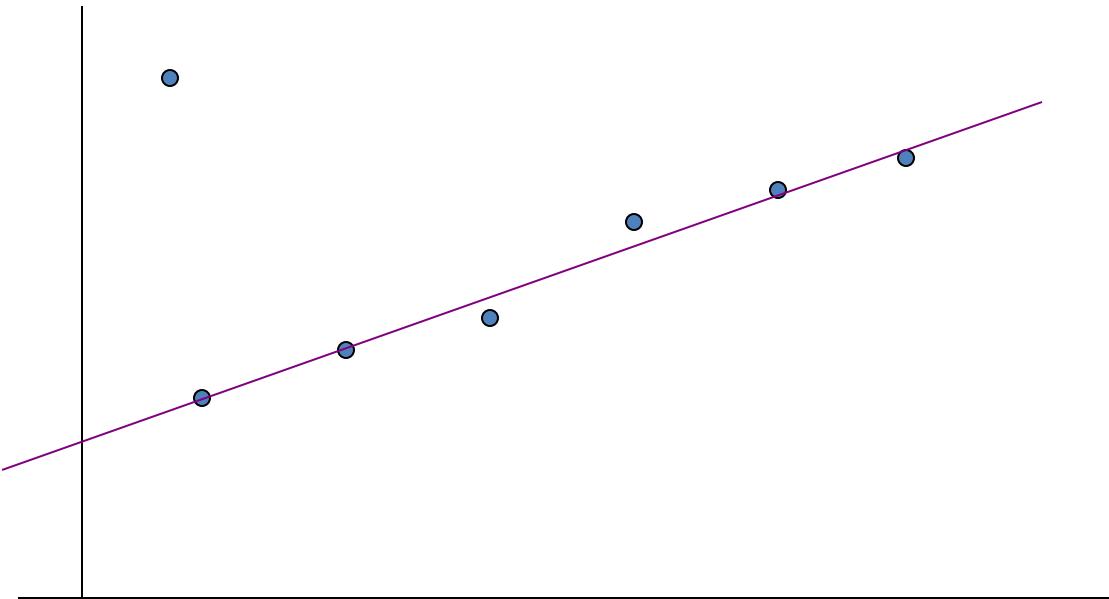
# Omnidirectional example



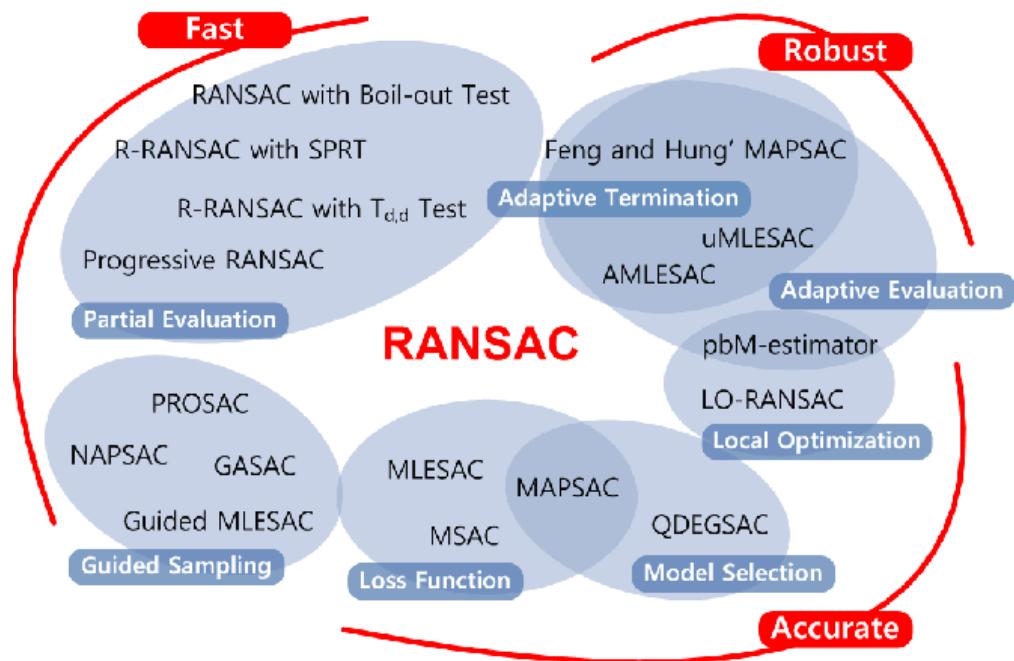
Images by Branislav Micusik, Tomas Pajdla,  
[cmp.felk.cvut.cz/ demos/Fishepi/](http://cmp.felk.cvut.cz/demos/Fishepi/)

# Simpler Example

- Fitting a straight line

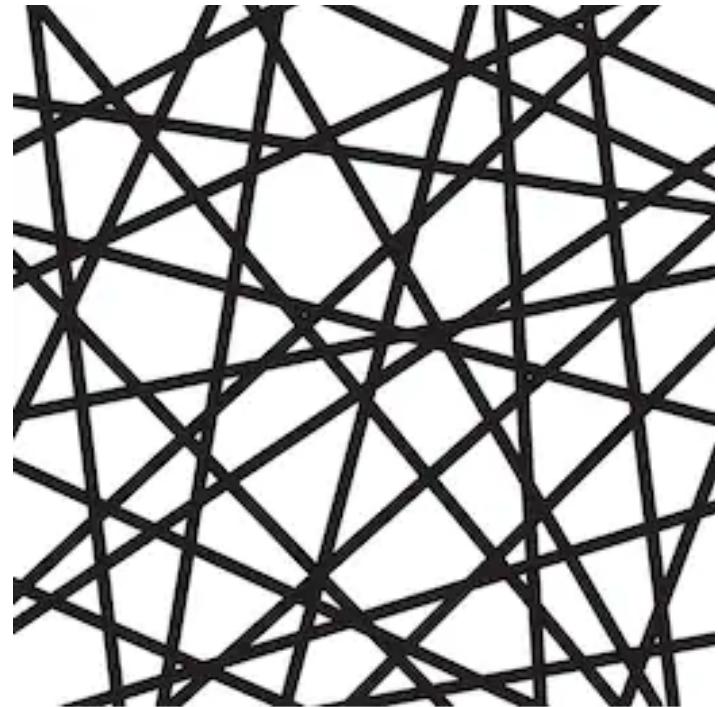


# Discard Outliers



- No point with  $d > t$
- RANSAC:
  - RANdom SAmple Consensus
  - Fischler & Bolles 1981
  - Copes with a large proportion of outliers

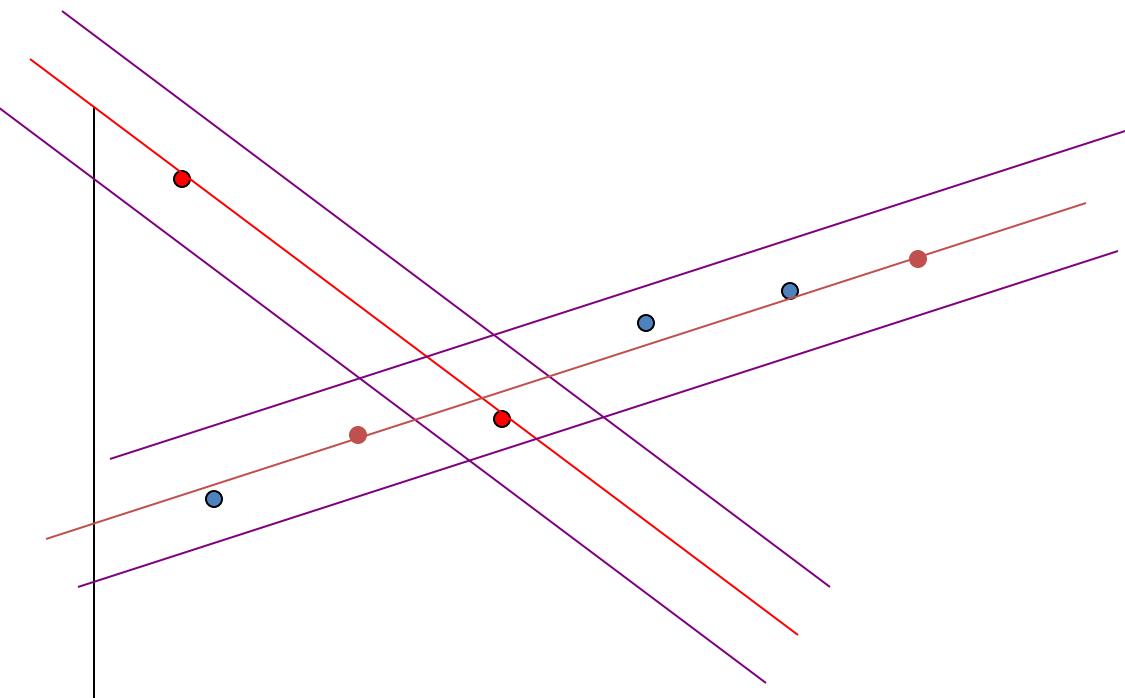
# Main Idea



shutterstock.com • 547881814

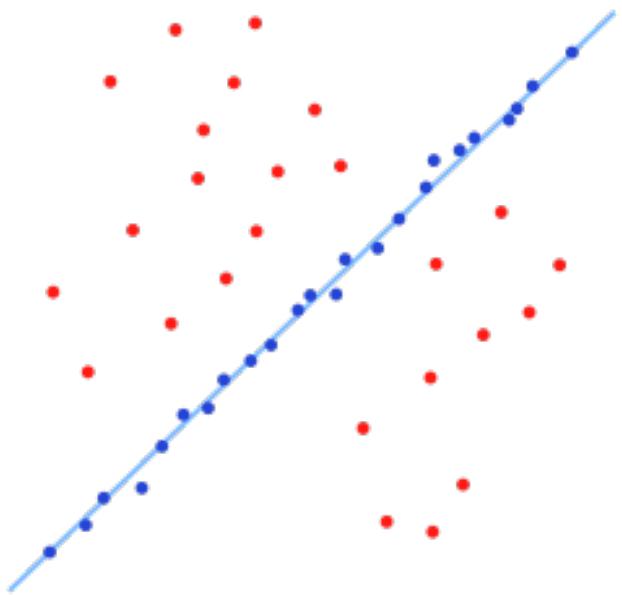
- Select 2 points at random
- Fit a line
- “Support” = number of inliers
- Line with most inliers wins

# Why will this work ?



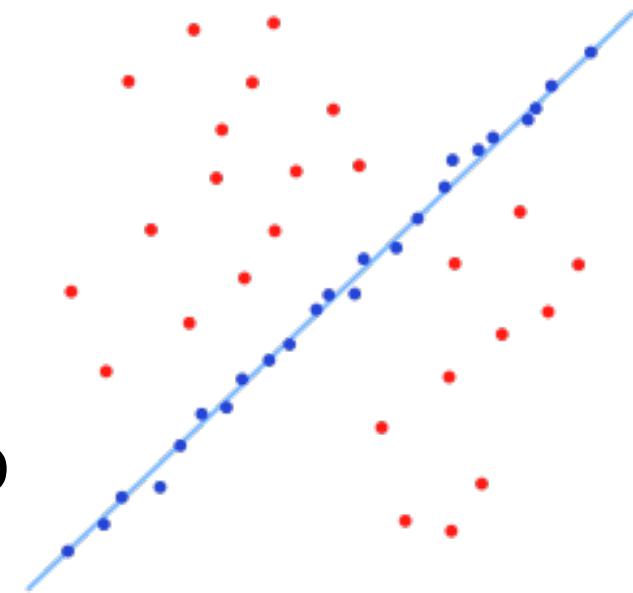
# Best Line has most support

- More support -> better fit



# RANSAC

- Objective:
  - Robust fit of a model to data  $D$
- Algorithm
  - Randomly select  $s$  points
  - Instantiate a model
  - Get consensus set  $D_i$
  - If  $|D_i| > T$ , terminate and return model
  - Repeat for  $N$  trials, return model with max  $|D_i|$



# In General



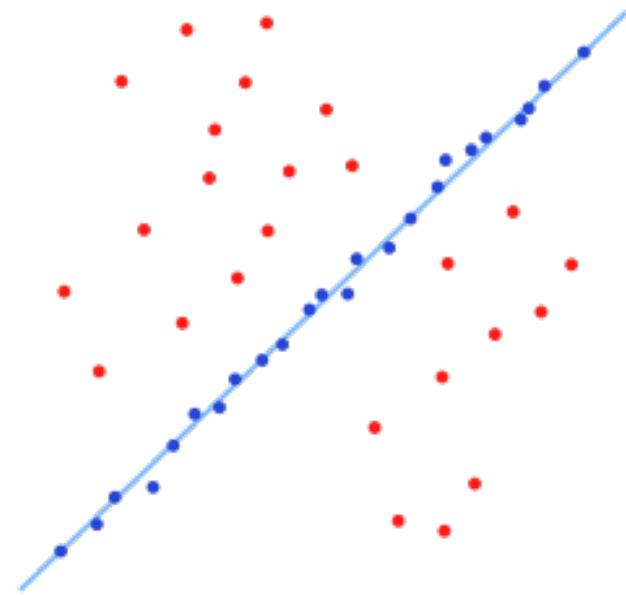
- Fit a more general model
- Sample = minimal subset
  - Translation ?
  - Homography ?
  - Euclidean transform ?

# Example



- Euclidean: needs 2 correspondences ( $2*2 >= 3$ )
- Here correct hypothesis has support of 4 (out of 5)
- Including red into minimal sample (of 2) would **likely** yield low support

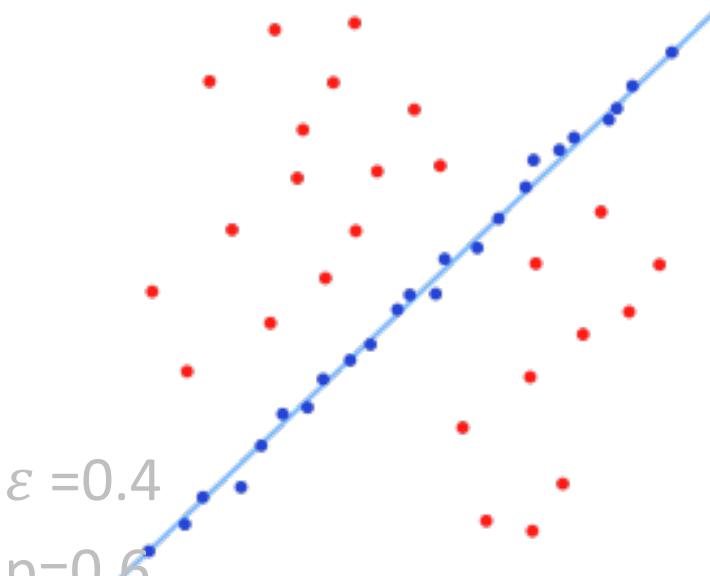
# How many samples ?



- We want: at least one sample with **all inliers**
- Can't guarantee: probability  $P$
- E.g.  $P = 0.99$

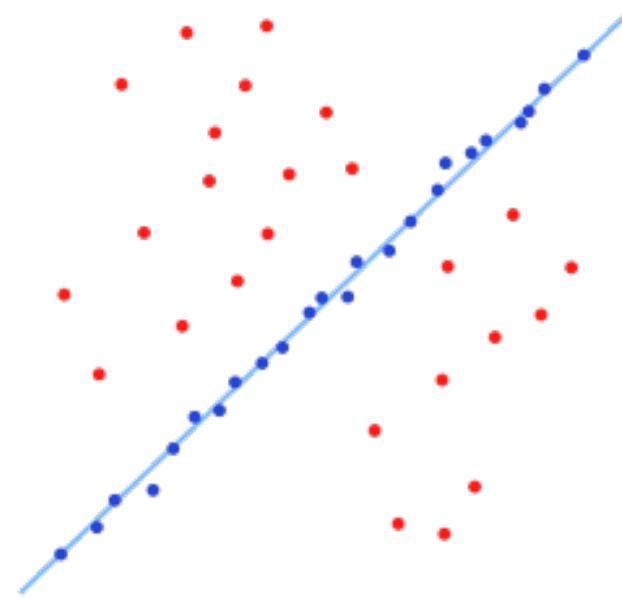
# Calculate N

- If  $\varepsilon$  = outlier probability
  - proportion of inliers  $p = 1 - \varepsilon$
  - $P(\text{sample with all inliers}) = p^s$
  - $P(\text{sample with an outlier}) = 1 - p^s$
  - $P(N \text{ samples an outlier}) = (1 - p^s)^N$
  - We want  $P(N \text{ samples an outlier}) < 1 - P$  (e.g. 0.01)
  - $(1 - p^s)^N < 1 - P$
  - $N > \log(1 - P) / \log(1 - p^s)$
- $\varepsilon = 0.4$   
 $p = 0.6$   
 $s = 2 \rightarrow p^s = 0.36$   
0.64  
 $N = 3 \rightarrow 0.26$   
 $0.64^N < 0.01$   
 $N > 10.3$

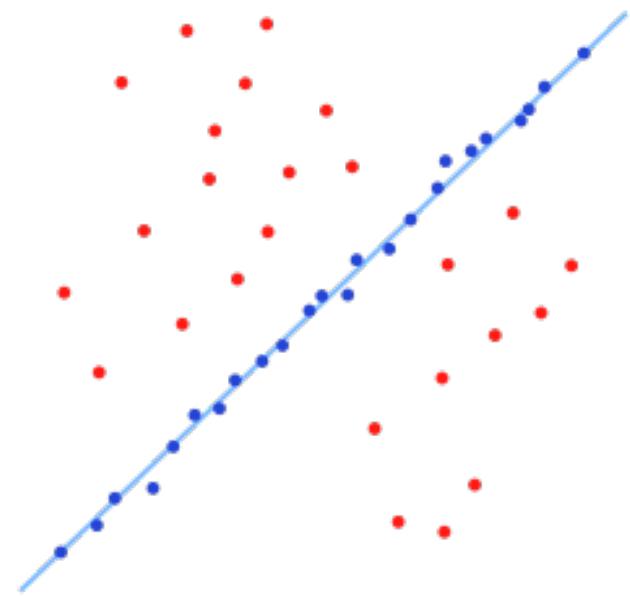


# Example

- $P=0.99$
- $s=2$ 
  - $\varepsilon = 5\%$   $\Rightarrow N=2$
  - $\varepsilon = 50\%$   $\Rightarrow N=17$
- $s=4$ 
  - $\varepsilon = 5\%$   $\Rightarrow N=3$
  - $\varepsilon = 50\%$   $\Rightarrow N=72$
- $s=8$ 
  - $\varepsilon = 5\%$   $\Rightarrow N=5$
  - $\varepsilon = 50\%$   $\Rightarrow N=1177$



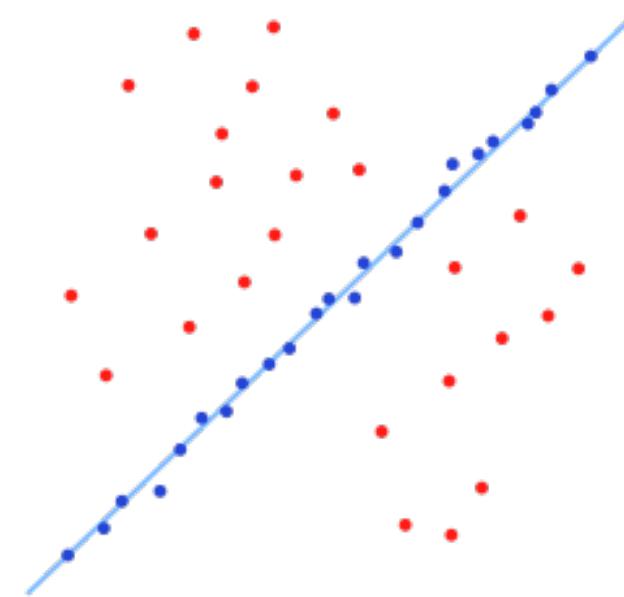
# Remarks



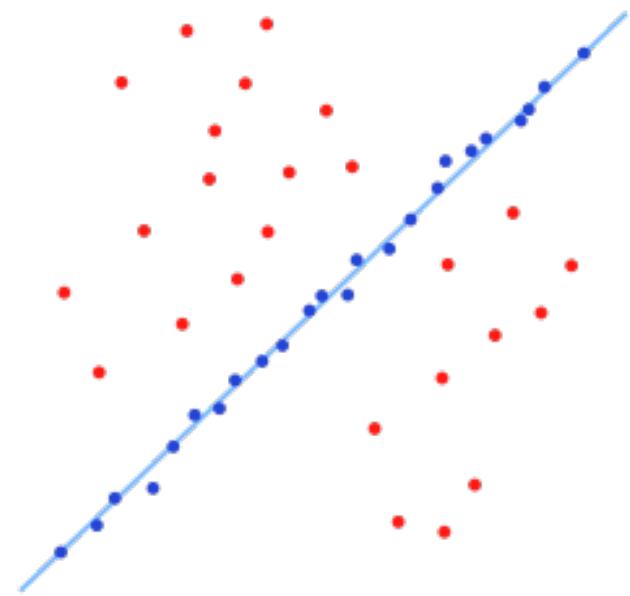
- $N = f(\varepsilon)$ , **not** the number of points
- $N$  increases steeply with  $s$

# Distance Threshold

- Requires noise distribution
- Gaussian noise with  $\sigma$
- Chi-squared distribution with DOF m
  - 95% cumulative:
  - Line, F:  $m=1, t=3.84 \sigma^2$
  - Translation, homography:  $m=2, t=5.99 \sigma^2$
- I.e.  $\rightarrow$  95% prob that  $d < t$  is inlier



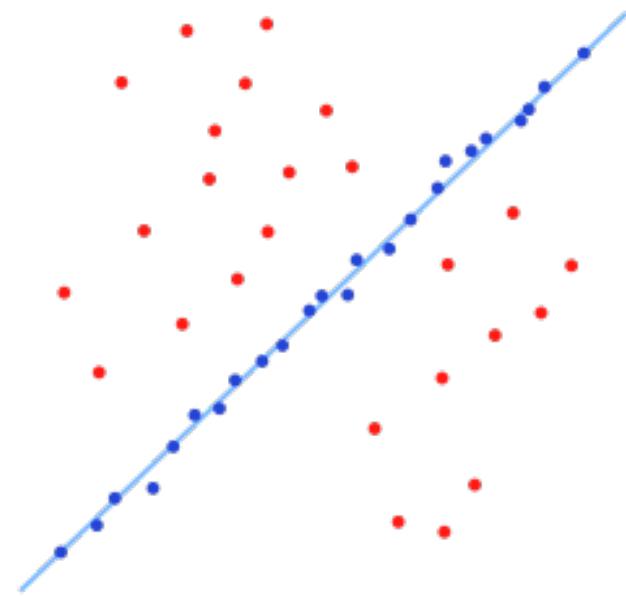
# Threshold T



- Terminate if  $|D_i| > T$
- Rule of thumb:  $T \approx \# \text{inliers}$
- So,  $T = (1 - \varepsilon)n = pn$

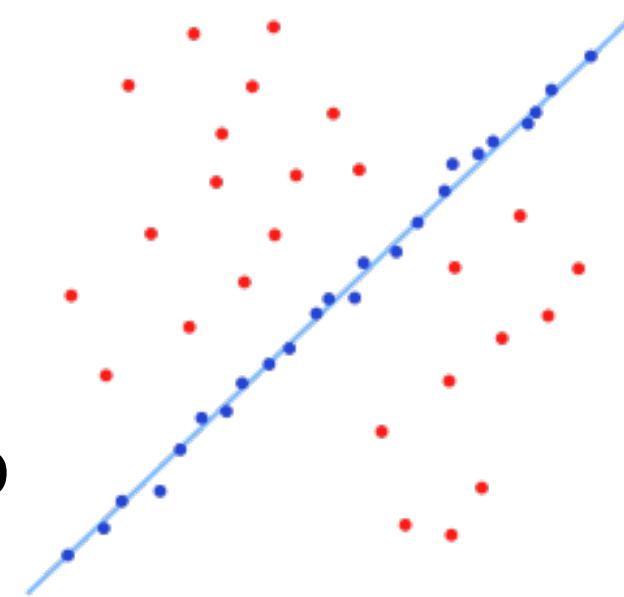
# Adaptive N

- When  $\varepsilon$  is unknown ?
- Start with  $\varepsilon = 50\%$ ,  $N=\infty$
- Repeat:
  - Sample  $s$ , fit model
  - $\rightarrow$  update  $\varepsilon$  as  $|\text{outliers}|/n$
  - $\rightarrow$  set  $N=f(\varepsilon, s, p)$
- Terminate when  $N$  samples seen

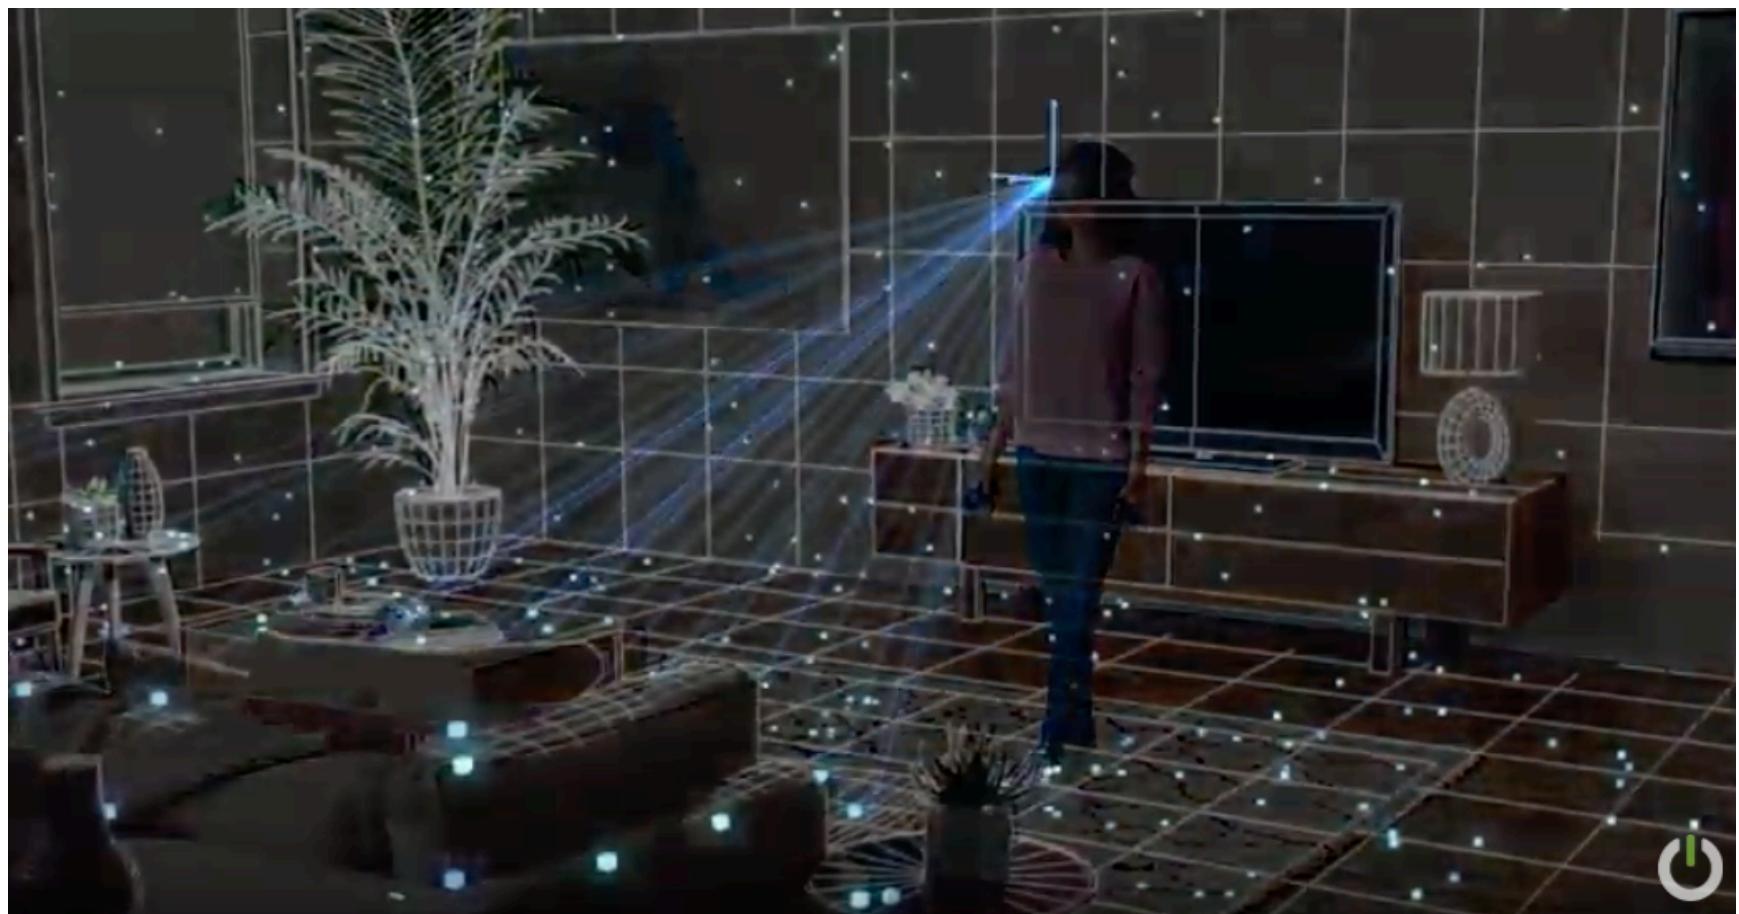


# Summary: RANSAC

- Objective:
  - Robust fit of a model to data  $D$
- Algorithm
  - Randomly select  $s$  points
  - Instantiate a model
  - Get consensus set  $D_i$
  - If  $|D_i| > T$ , terminate and return model
  - Repeat for  $N$  trials, return model with max  $|D_i|$



# Pose Estimation in VR



<https://youtu.be/nrj3JE-NHMw>

Frank Dellaert Fall 2020

# Review: 2D Alignment



- Input:
  - A set of matches  $\{(x_i, x'_i)\}$
  - A parametric model  $f(x; p)$
- Output:
  - Best model  $p^*$
- How?

# Now: 3D-2D Alignment



- Input:
  - A set of 3D->2D matches  $\{(X_i, x_i)\}$
  - A parametric model  $f(X; p)$
- Output:
  - Best model  $p^*$
- How?

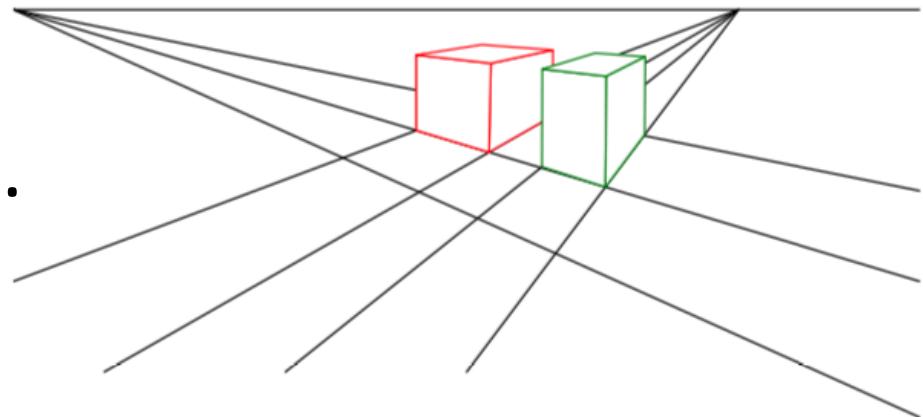
# Pose Estimation



- Input:
  - A set of 2D measurements  $x_i$  of known 3D points  $3D X_i$
  - Parametric model is camera matrix  $P$ , i.e.,  $x = f(X; P)$
- Output:
  - Best camera matrix  $P$
- How?

# Review: Projective Camera Matrix

- Chapter 2 in book
- Homogeneous coord.
- 3D TO 2D projection:



$$\mathbf{x} = \mathbf{K}[\mathbf{R}|\mathbf{t}]\mathbf{X} = \mathbf{P}\mathbf{X}$$

where  $P = 3 \times 4$  camera matrix

and  $K$  the  $3 \times 3$  calibration

$$\mathbf{K} = \begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

# Camera Extrinsics: a Pose in 3D

- What is the geometric meaning of  $R$  and  $t$  ??
- Intuitive: camera is at a position  ${}_w t_c$   
Indices say: camera *in* world coordinate frame



$${}_w t_c$$

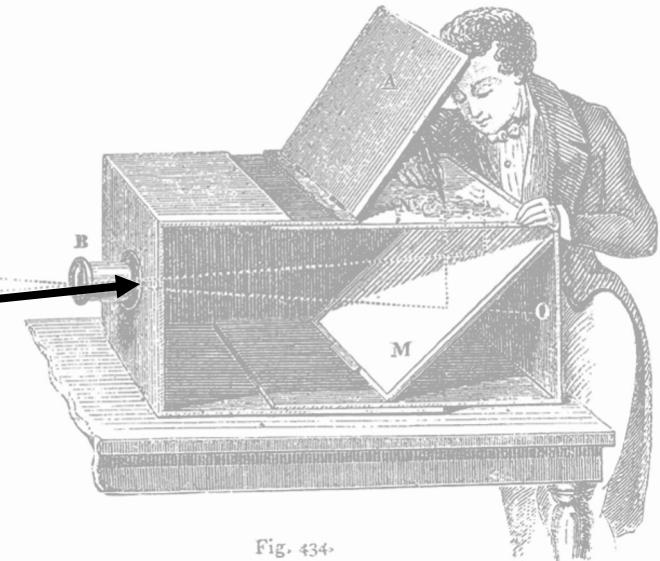
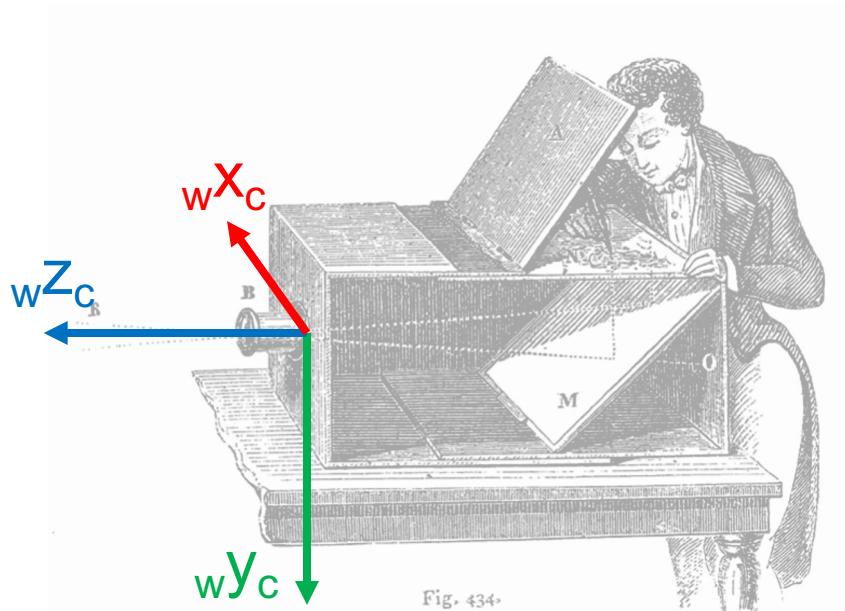


Fig. 434.

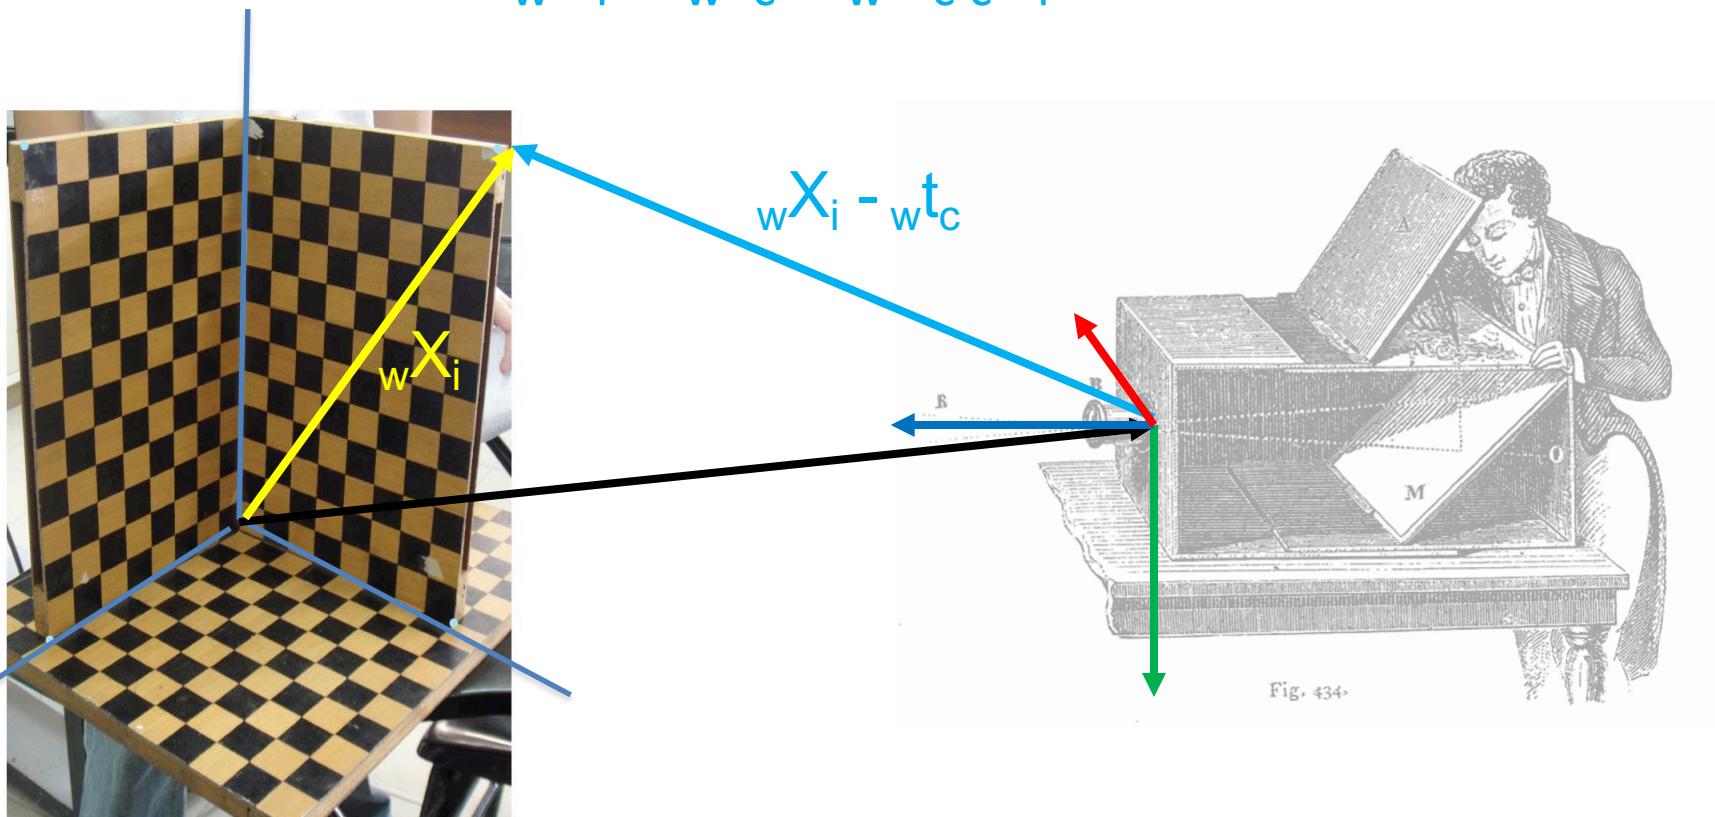
# Camera Extrinsic: a Pose in 3D

- What is the geometric meaning of  $R$  and  $t$  ??
- Rotation is given by 3x3 matrix  ${}^w R_c$  whose *columns* are the camera axes  ${}^w X_c$ ,  ${}^w Y_c$ ,  ${}^w Z_c$



# Camera Extrinsics: a Pose in 3D

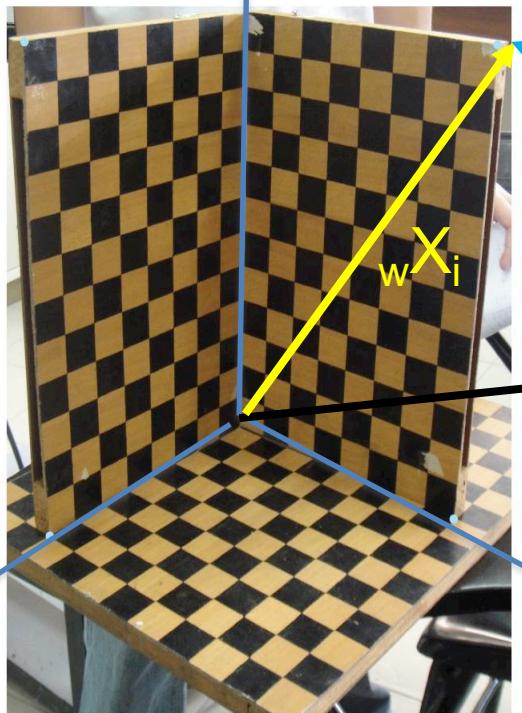
- What is the geometric meaning of R and t ??
- Transforming point  $X_i$  from world to camera coordinates:  ${}_{\text{w}}X_i - {}_{\text{w}}t_c = {}_{\text{w}}R_{\text{c}\text{c}}{}_{\text{c}}X_i$



# Camera Extrinsics: a Pose in 3D

- Expressed in homogeneous coordinates:

$$\begin{aligned} {}_c X_i &= {}_w R_c^T ({}_w X_i - {}_w t_c) = {}_w R_c^T [I | - {}_w t_c] {}_w X_i \\ &= [{}_w R_c^T I | - {}_w R_c^T {}_w t_c] {}_w X_i \\ &= [{}_c R_w | {}_c t_w] {}_w X_i \\ &= [R | t] {}_w X_i \end{aligned}$$



${}_c X_i$

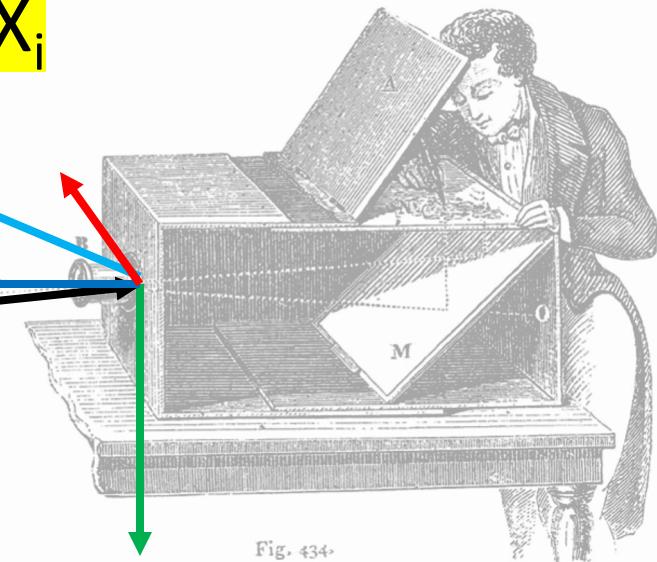


Fig. 434.

# Camera Extrinsic: a Pose in 3D

- Conclusion: when people write  ${}_{\text{c}}X_i = [R|t] {}_{\text{w}}X_i$  they are talking about (unintuitive)  $[{}_{\text{c}}R_w | {}_{\text{c}}t_w]$
- We like use (intuitive)  ${}_{\text{c}}X_i = {}_{\text{w}}R_c^T [I| - {}_{\text{w}}t_c] {}_{\text{w}}X_i$

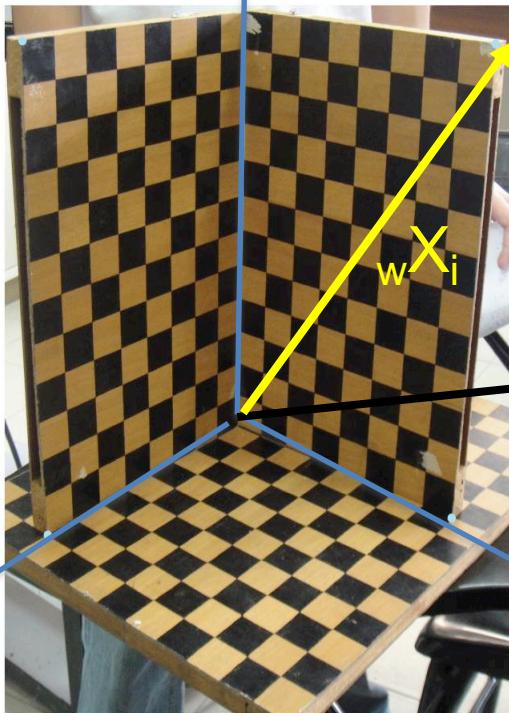
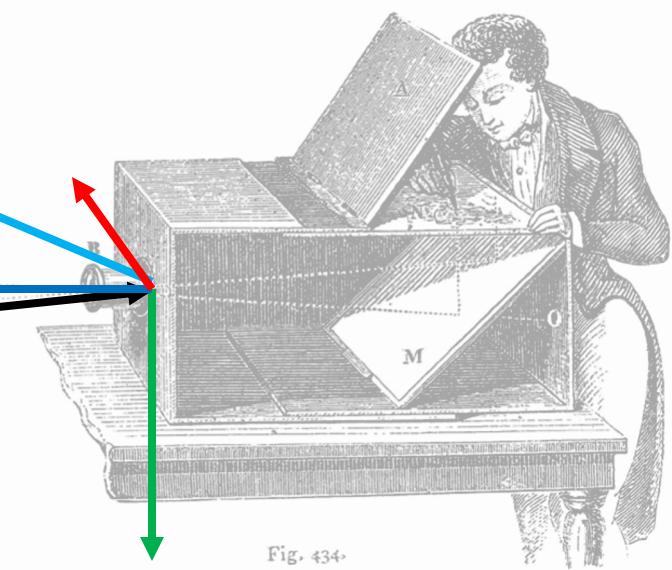
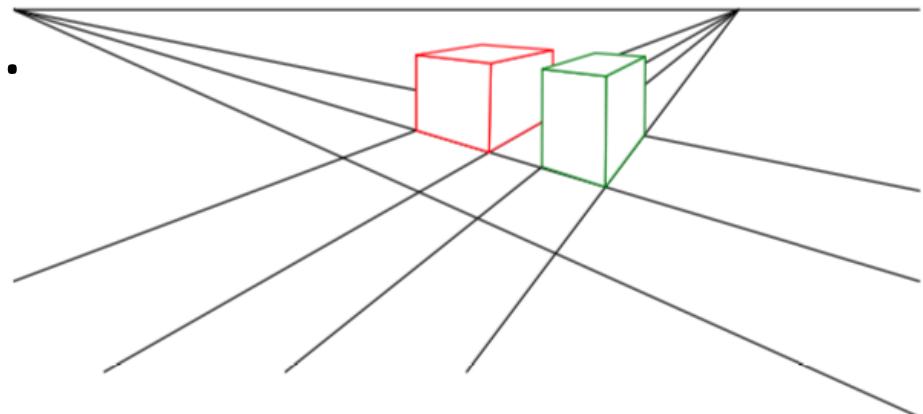
 $cX_i$  $R$ 

Fig. 434.

# Revision: Projective Camera Matrix

- Homogeneous coord.
- 3D TO 2D projection:



Camera-centric:  $\mathbf{x} = \mathbf{K}_{\mathbf{c}} \mathbf{R}_{\mathbf{w}} | {}_{\mathbf{c}} \mathbf{t}_{\mathbf{w}} ] \mathbf{X} = \mathbf{P} \mathbf{X}$

World-centric:  $\mathbf{x} = \mathbf{K}_{\mathbf{w}} \mathbf{R}_{\mathbf{c}}^T [ \mathbf{I} | - {}_{\mathbf{w}} \mathbf{t}_{\mathbf{c}} ] \mathbf{X} = \mathbf{P} \mathbf{X}$

$P$  = same  $3 \times 4$  camera matrix

and  $K$  the  $3 \times 3$  calibration

$$\mathbf{K} = \begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

# Looking at the (opaque) camera matrix

Can you interpret the columns of  $P$  with entities in the scene?

$$P = [P^1 \ P^2 \ P^3 \ P^4]$$

Answer:

$P^1$  == the image of  $[1 \ 0 \ 0 \ 0]$

$P^2$  == the image of  $[0 \ 1 \ 0 \ 0]$

$P^3$  == the image of  $[0 \ 0 \ 1 \ 0]$

$P^4$  == the image of  $[0 \ 0 \ 0 \ 1]$

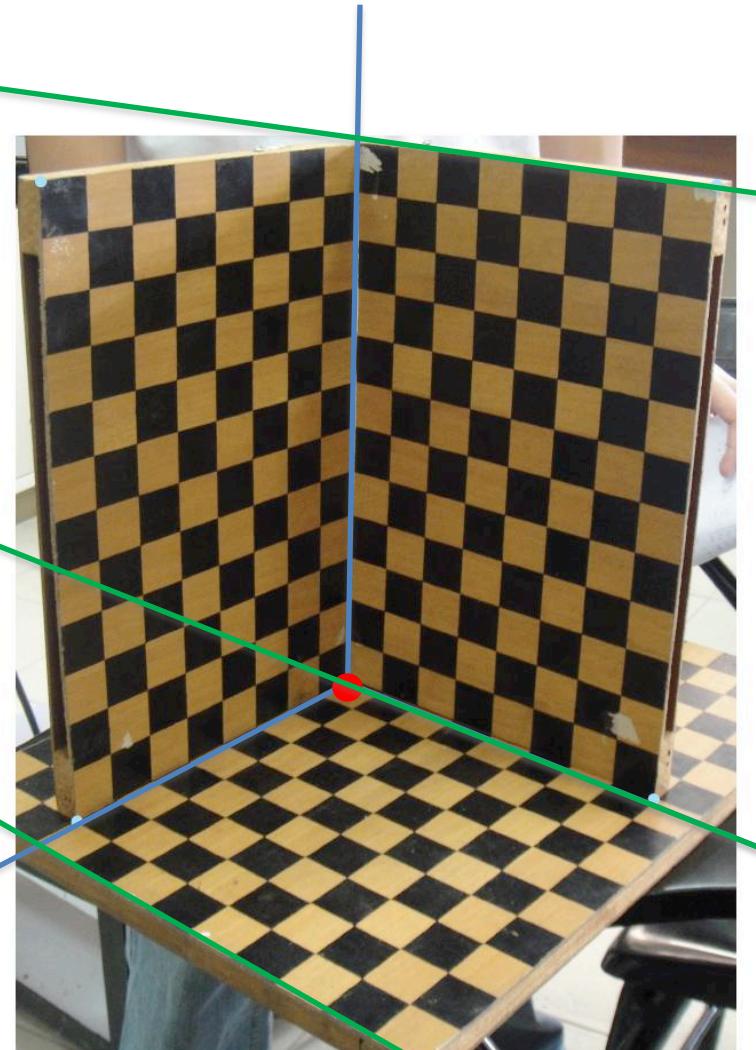
What are those ?

$[0 \ 0 \ 0 \ 1]$  is easy...

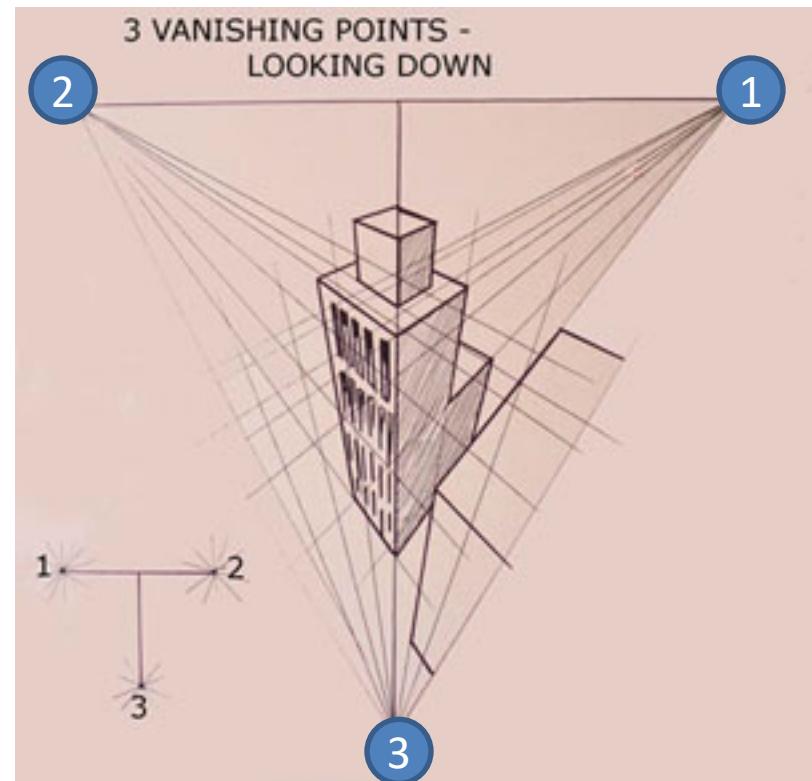
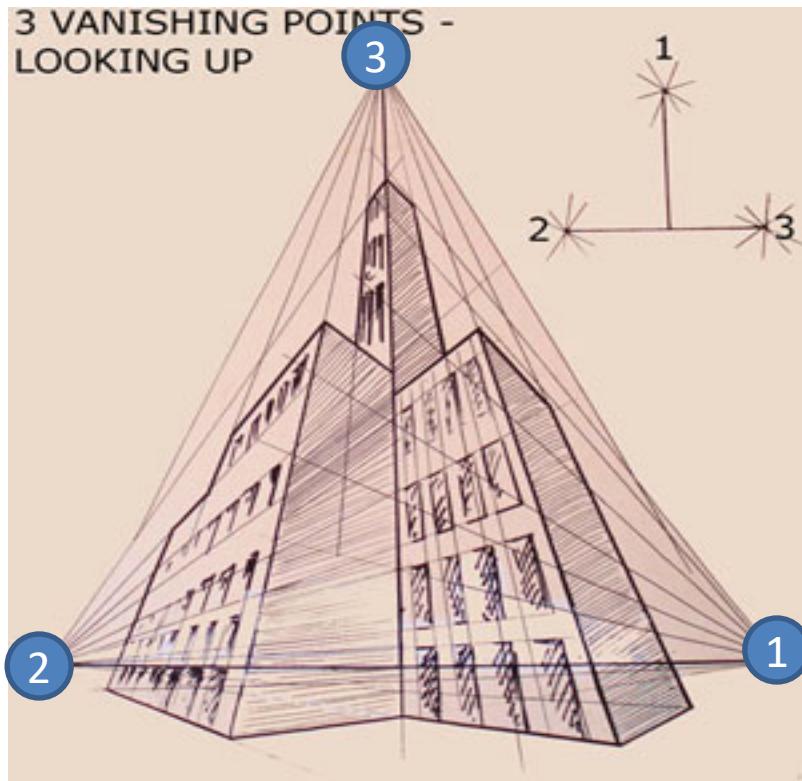
Answer:

$[0 \ 0 \ 0 \ 1]$  is the origin, so  $P^4$  is the image of the origin.

$[1 \ 0 \ 0 \ 0]$  is a point at infinity in the X-direction, so it is the vanishing point of all lines parallel with the X direction!



# Vanishing points, revisited



Columns of P !

$$P = [P^1 \quad P^2 \quad P^3 \quad P^4]$$

$P^4$  is arbitrary: wherever you defined the world origin.

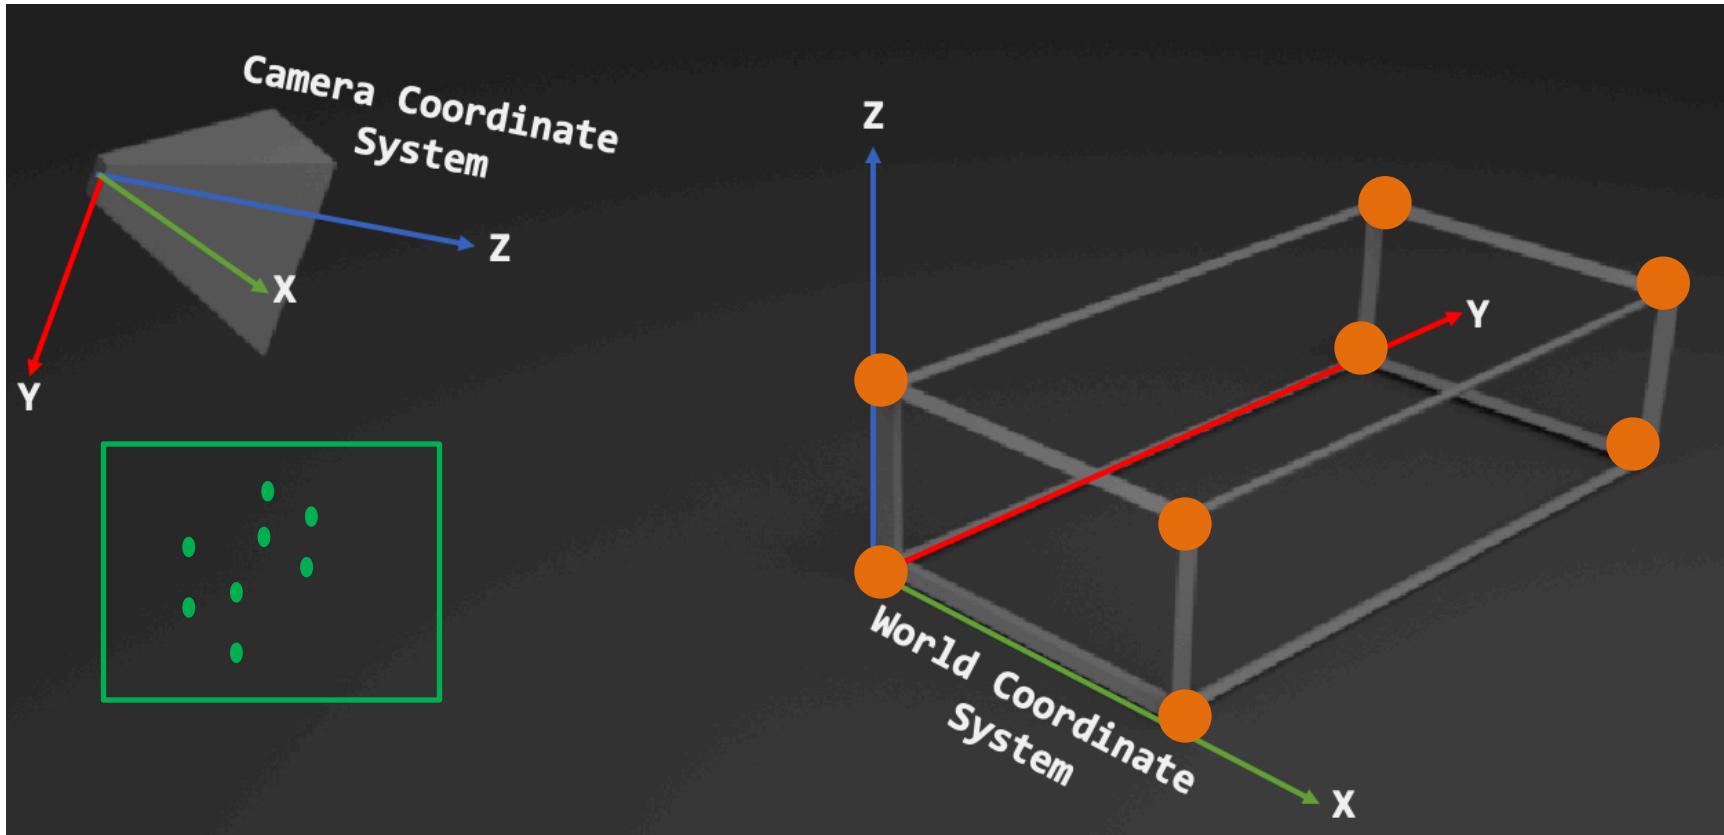
<https://www.artinstructionblog.com/perspective-drawing-tutorial-for-artists-part-2>

Frank Dellaert Fall 2020

# Back to Pose Estimation!

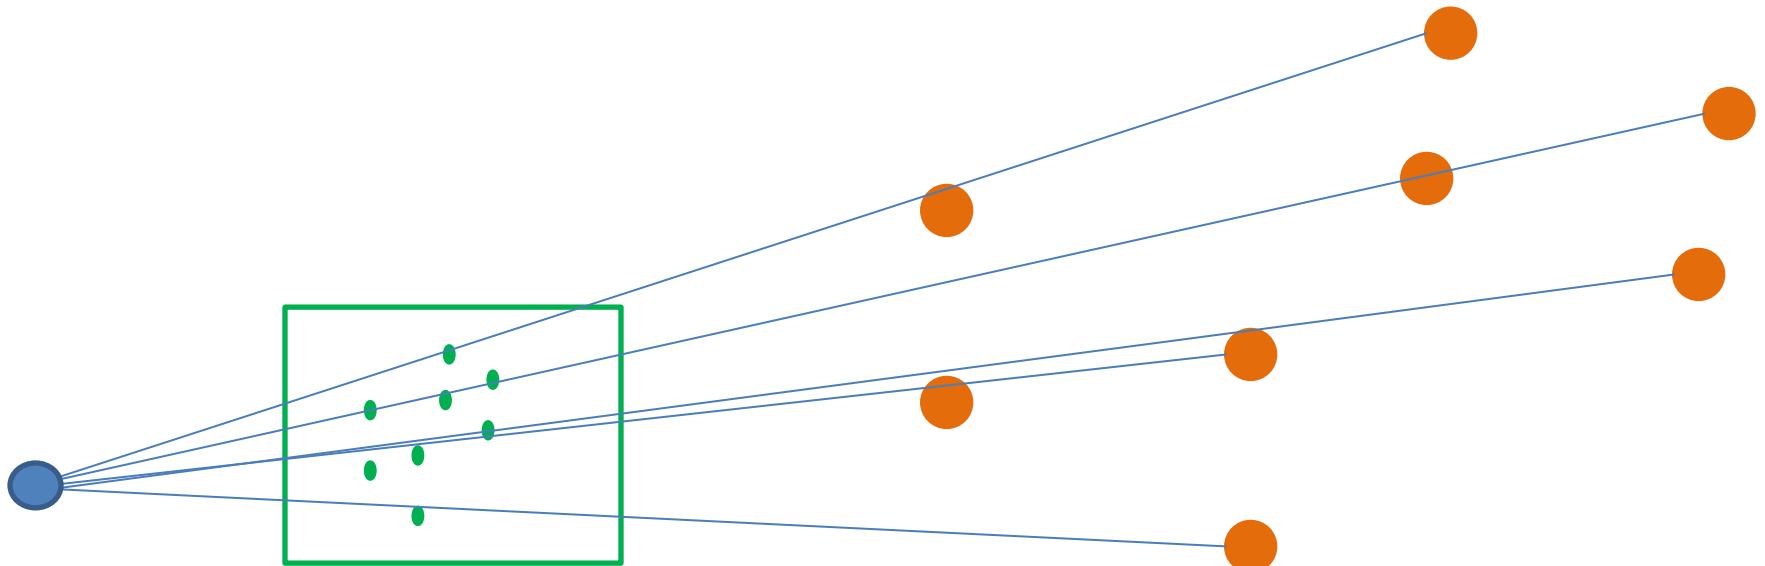
- Simple algorithm: just measure the coordinates of the origin and the three vanishing points?
- Does not work ☹:
  - Columns are only measured up to a scale.
  - $4 \text{ points} * 2\text{DOF} = \text{only } 8 \text{ DOF!}$  Missing  $11-8=3$
  - 3 missing numbers are exactly those scales.

# Least Squares Pose Estimation...



- Input:
  - A set of **2D measurements**  $x_i$  of known 3D **points3D**  $X_i$
  - Parametric model is camera matrix  $P$ , i.e.,  $x = f(X; P)$
- Output:
  - Best camera matrix  $P$

# Pose estimation = “Resectioning”



$$\mathbf{x} = f(\mathbf{X}_w; \mathbf{P}) = \mathbf{P}\mathbf{X}_w = \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \cong \begin{bmatrix} s \cdot u \\ s \cdot v \\ s \end{bmatrix} = \begin{bmatrix} p_{11} & p_{12} & p_{13} & p_{14} \\ p_{21} & p_{22} & p_{23} & p_{24} \\ p_{31} & p_{32} & p_{33} & p_{34} \end{bmatrix} \begin{bmatrix} x_w \\ y_w \\ z_w \\ 1 \end{bmatrix}$$

$$\arg \min_{\hat{\mathbf{P}}} \sum_{i=1}^N \|\hat{\mathbf{P}}\mathbf{X}_w^i - \mathbf{x}^i\|_2.$$

- Opposite of triangulation.

# Pose estimation

$$\arg \min_{\hat{\mathbf{P}}} \sum_{i=1}^N \|\hat{\mathbf{P}} \mathbf{X}_w^i - \mathbf{x}^i\|_2.$$

- In project 4, you will use `scipy.optimize.least_squares` to do exactly that. Working knowledge of 3D poses will be required.
- Note before we compute the 2D reprojection error we need to convert back  $PX$  to non-homogeneous coordinates:

$$x_i = \frac{p_{00}X_i + p_{01}Y_i + p_{02}Z_i + p_{03}}{p_{20}X_i + p_{21}Y_i + p_{22}Z_i + p_{23}}$$
$$y_i = \frac{p_{10}X_i + p_{11}Y_i + p_{12}Z_i + p_{13}}{p_{20}X_i + p_{21}Y_i + p_{22}Z_i + p_{23}}$$

[https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.least\\_squares.html](https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.least_squares.html)